

The project for the machine learning course has officially launched. This year, we have organized a Kaggle competition in which you will all compete against each other. The goal is to build a classifier that is able to distinguish between twelve kinds of animals. You have two months time to tinker away at your models and become the best model builder the world has ever seen. Just like a real Kaggle competition, the students who hold the top three submissions at the end of the competition on the private leaderboard receive eternal fame and glory, but also a prize in the form of a copy of the book *'Human Compatible: AI and the Problem of Control'*, by Stuart Russell. For now, the competition is restricted to non-neural network solutions; a smaller separate competition will be launched later where deep learning approaches are encouraged.

We require you to write a report on your solution(s), and we will first and foremost grade you on that report, so do start writing it early on! Besides your final report at the end of the competition, we expect you to hand in a smaller intermediate report as well, so that we can track everyone's progress and give useful feedback during the competition. More information on the reports can be found on the Kaggle competition page.

The deadline for the competition is **Friday, January 15, 23:59**. The deadline for the report is **Sunday, January 17, 23:59**. The deadline for the intermediate report is **Friday, December 11, 23:59**. Start early, because you will probably encounter many hurdles along the way, and fine-tuning models, analyzing your results and then doing this countless times over and over again simply takes time. Your reports and code should be uploaded to a GitHub classroom, which we'll create soon.

Detailed information can be found on the Kaggle competition site, and I'll be covering the challenge during the WPO lectures next week, so try to read and set up everything (and prepare your questions) by then. Definitely cover everything on the Kaggle page: from the description, to the rules, to what is written about the dataset. The dataset comes with a few useful notebooks and helper scripts to get you started, so after reading the competition page, we would advise you to start there. The link to join the competition is [here \(Koppelingen naar een externe site.\)](#). Do not share this link outside of the class.

Updates will be posted on Canvas, and if anything changes to the Kaggle competition site (e.g. a new version of the dataset or a change in evaluation metric), there will be instructions and a notification here. If anything is still unclear, or if you have questions, do not hesitate to mail me at andries.rosseau@vub.be.